

```
In [2]: # Task 1: ML Basics - Iris Dataset
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, classification_r
```

```
In [3]: # Load Dataset
iris = load_iris()
X = iris.data
y = iris.target
```

```
In [4]: # Convert to DataFrame for analysis
df = pd.DataFrame(X, columns=iris.feature_names)
df['target'] = y
```


```
In [5]: # Map target labels to names
df['target_name'] = df['target'].map(dict(enumerate(iris.target_names)))
```

```
In [6]: # Display top records
print(df.head())
```


	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	target	target_name
0	0	setosa
1	0	setosa
2	0	setosa
3	0	setosa
4	0	setosa

```
In [7]: # Basic EDA
sns.pairplot(df, hue='target_name')
plt.title("Pairplot of Iris Features")
plt.show()
```



```
In [8]: # Feature correlation heatmap
plt.figure(figsize=(8,6))
sns.heatmap(df.drop(columns='target_name').corr(), annot=True, cmap='coolwarm')
plt.title("Feature Correlation")
plt.show()
```



```
In [9]: # Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [10]: # Feature Scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [11]: # ----- Model 1: Logistic Regression -----
log_model = LogisticRegression(max_iter=200)
log_model.fit(X_train_scaled, y_train)
y_pred_log = log_model.predict(X_test_scaled)
```

```
In [12]: # Evaluation
print("❖ Logistic Regression")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_log))
print("Accuracy:", accuracy_score(y_test, y_pred_log))
print("Classification Report:\n", classification_report(y_test, y_pred_log))
```

```

◇ Logistic Regression
Confusion Matrix:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
Accuracy: 1.0
Classification Report:
precision    recall   f1-score   support
0            1.00    1.00     1.00      10
1            1.00    1.00     1.00       9
2            1.00    1.00     1.00      11

accuracy                   1.00      30
macro avg                 1.00     1.00      1.00      30
weighted avg               1.00     1.00     1.00      30

```

```
In [13]: # ----- Model 2: K-Nearest Neighbors -----
knn_model = KNeighborsClassifier(n_neighbors=3)
knn_model.fit(X_train_scaled, y_train)
y_pred_knn = knn_model.predict(X_test_scaled)
```

```
In [14]: # Evaluation
print("\n◇ K-Nearest Neighbors (k=3)")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_knn))
print("Accuracy:", accuracy_score(y_test, y_pred_knn))
print("Classification Report:\n", classification_report(y_test, y_pred_knn))
```

```

◇ K-Nearest Neighbors (k=3)
Confusion Matrix:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
Accuracy: 1.0
Classification Report:
precision    recall   f1-score   support
0            1.00    1.00     1.00      10
1            1.00    1.00     1.00       9
2            1.00    1.00     1.00      11

accuracy                   1.00      30
macro avg                 1.00     1.00      1.00      30
weighted avg               1.00     1.00     1.00      30

```

```
In [ ]:
```

Task 1: Introduction to Machine Learning with Scikit-learn

Dataset Used: Iris Dataset

Objective: Understand core ML concepts and apply classification algorithms to a real-world dataset.

Dataset Overview

The Iris dataset consists of 150 samples of iris flowers from three species: Setosa, Versicolor, and Virginica. Each sample has four features:

- Sepal length
 - Sepal width
 - Petal length
 - Petal width
-

Model Selection

We applied **two supervised classification algorithms**:

1. Logistic Regression

- Chosen as a simple and efficient baseline model.
- Suitable for multi-class classification using the one-vs-rest approach.
- Performs well on linearly separable data like Iris.

2. K-Nearest Neighbors (KNN)

- Chosen for its simplicity and non-parametric nature.
 - Works based on distance similarity, making it good for smaller datasets.
-

Model Evaluation

Both models were trained on 80% of the dataset and tested on the remaining 20%. Feature scaling was applied using `StandardScaler`.

Model	Accuracy	Confusion Matrix	Comments
Logistic Regression	100%	No misclassifications	Linear model, fast & interpretable
KNN (k=3)	100%	No misclassifications	Performs well when data is clean and low-dimensional

Evaluation metrics used:

- **Confusion Matrix:** Showed perfect classification for all classes.
 - **Accuracy Score:** 1.0 for both models.
 - **Precision, Recall, F1-score:** All metrics were perfect (1.00).
-

Conclusion

Both models performed exceptionally well on the Iris dataset. In real-world settings, such performance might be rare due to noise and complexity in data. However, this task successfully demonstrated:

- Dataset analysis
- Model training and comparison
- Evaluation using industry-standard metrics

For future work, these models can be tested with:

- Noisy/real-world datasets
- Hyperparameter tuning (e.g., K in KNN)
- Cross-validation