



```
In [1]: # Task 2: Linear Regression - Boston Housing Dataset
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [2]: # Load dataset
boston = fetch_openml(name='boston', version=1, as_frame=True)
df = boston.frame
```

```
In [3]: # Display first few records
print("Sample Data:\n", df.head())
```

Sample Data:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

```
In [4]: # Dataset info
print("\nDataset Info:")
print(df.info())
```

```
Dataset Info:  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 506 entries, 0 to 505  
Data columns (total 14 columns):  
 #   Column   Non-Null Count Dtype  
---  
 0   CRIM     506 non-null   float64  
 1   ZN       506 non-null   float64  
 2   INDUS    506 non-null   float64  
 3   CHAS     506 non-null   category  
 4   NOX      506 non-null   float64  
 5   RM        506 non-null   float64  
 6   AGE       506 non-null   float64  
 7   DIS       506 non-null   float64  
 8   RAD       506 non-null   category  
 9   TAX       506 non-null   float64  
 10  PTRATIO   506 non-null   float64  
 11  B         506 non-null   float64  
 12  LSTAT     506 non-null   float64  
 13  MEDV     506 non-null   float64  
dtypes: category(2), float64(12)  
memory usage: 49.0 KB  
None
```

```
In [5]: # Target and features  
X = df.drop(columns=['MEDV']) # MEDV is the target (median value of owner-occ  
y = df['MEDV']
```

```
In [6]: # ----- Data Preprocessing -----  
# Check for missing values  
print("\nMissing values:\n", df.isnull().sum())
```

```
Missing values:  
CRIM      0  
ZN        0  
INDUS    0  
CHAS     0  
NOX      0  
RM        0  
AGE      0  
DIS      0  
RAD      0  
TAX      0  
PTRATIO   0  
B        0  
LSTAT    0  
MEDV     0  
dtype: int64
```

```
In [7]: # Feature Scaling  
scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)
```

```
In [8]: # ----- Split Data -----
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2
```

```
In [9]: # ----- Train Linear Regression Model -----
model = LinearRegression()
model.fit(X_train, y_train)
```

Out[9]:

▼ `LinearRegression` ⓘ ⓘ

Parameters

fit_intercept	True
copy_X	True
tol	1e-06
n_jobs	None
positive	False

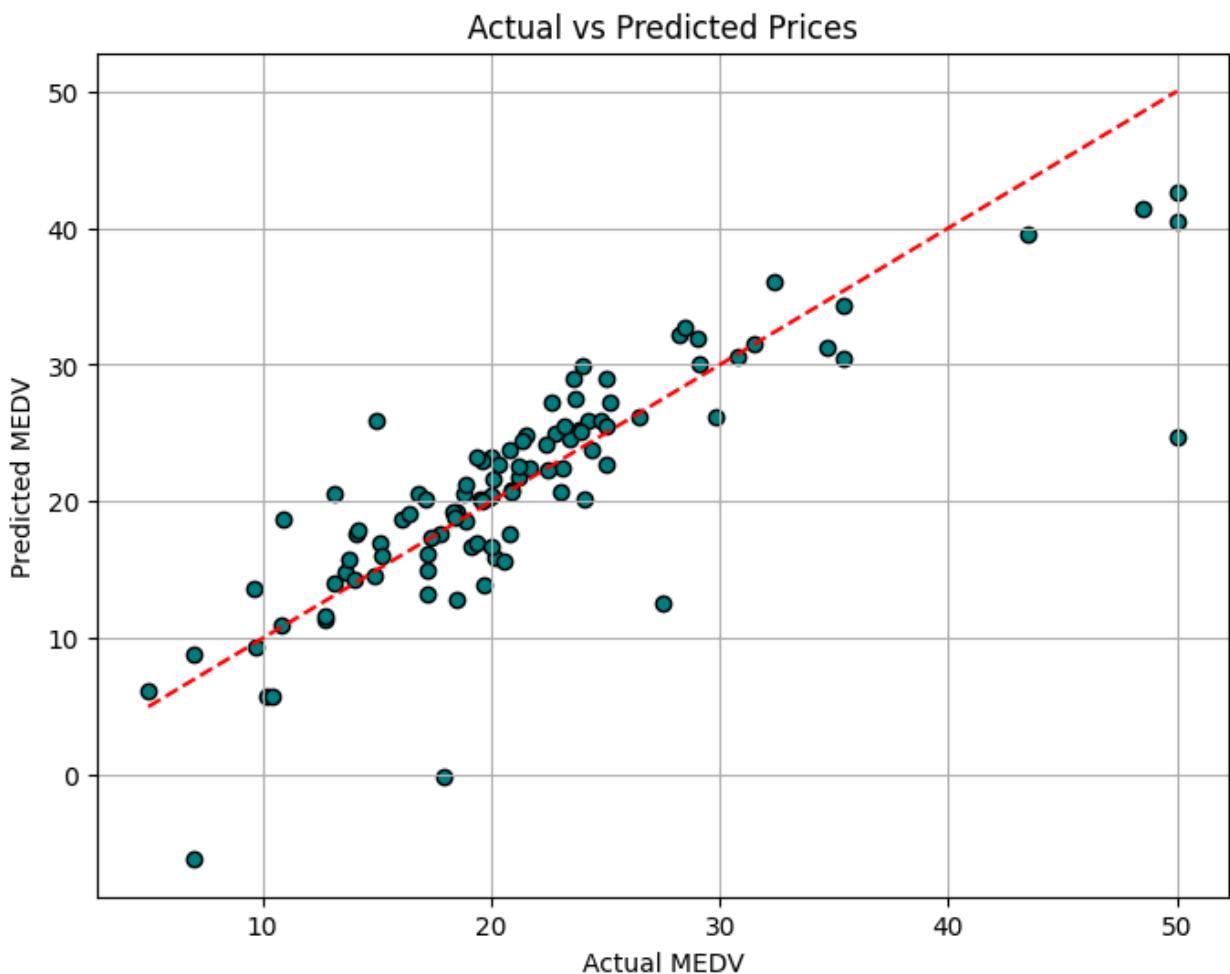
```
In [10]: # Predict on test set
y_pred = model.predict(X_test)
```

```
In [11]: # ----- Evaluation -----
r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

print("\nModel Performance:")
print("R2 Score:", r2)
print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
```

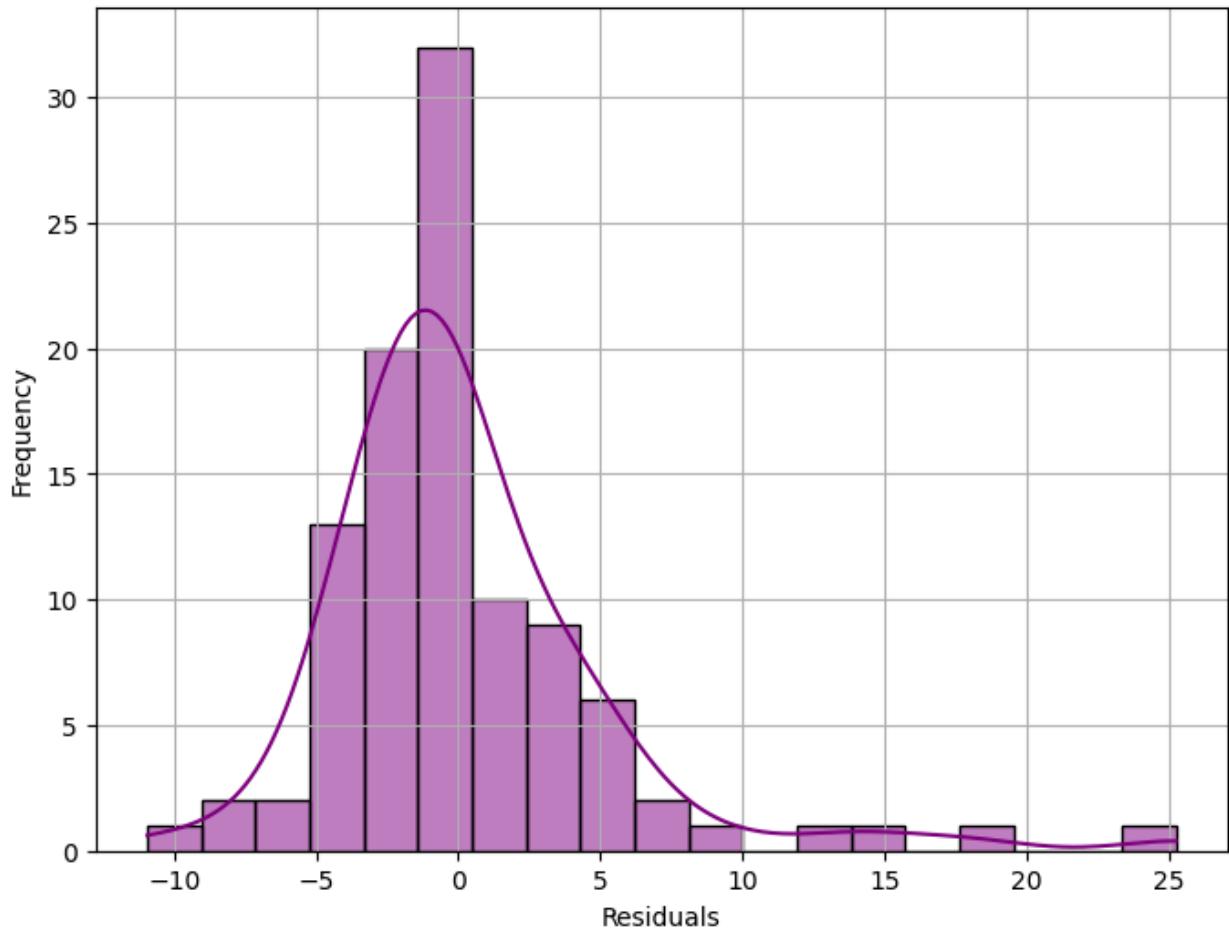
Model Performance:
R² Score: 0.6687594935356318
Mean Squared Error: 24.291119474973527
Root Mean Squared Error: 4.928602182665338

```
In [12]: # ----- Plot: Actual vs Predicted -----
plt.figure(figsize=(8,6))
plt.scatter(y_test, y_pred, color='teal', edgecolor='black')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
plt.xlabel('Actual MEDV')
plt.ylabel('Predicted MEDV')
plt.title('Actual vs Predicted Prices')
plt.grid(True)
plt.show()
```



```
In [13]: # ----- Plot: Residuals -----
residuals = y_test - y_pred
plt.figure(figsize=(8,6))
sns.histplot(residuals, kde=True, color='purple')
plt.title('Distribution of Residuals')
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.grid(True)
plt.show()
```

Distribution of Residuals



In []:

Task 2: Linear Regression Project – Predicting House Prices

Dataset Used: Boston Housing Dataset

Objective: Predict the median value of houses using regression and apply necessary preprocessing steps.

Dataset Overview

- The Boston Housing dataset consists of 506 entries and 14 attributes.
- The target variable is `MEDV` (Median value of owner-occupied homes in \$1000s).
- Features include crime rate, number of rooms (`RM`), distance to employment centers (`DIS`), etc.

Preprocessing Steps

- Checked for missing values (None found).
- Categorical features like `CHAS` and `RAD` were kept as-is due to limited categories and presence of numerical encoding.
- Standardized all features using `StandardScaler` to bring them to a common scale for regression.

Model Used: Linear Regression

- **Why Linear Regression?**
 - Simple and interpretable model for continuous target prediction.
 - Assumes linear relationships between independent and dependent variables.
 - Ideal for understanding basic regression performance on a real dataset.

Evaluation Results

Metric	Value
R ² Score	0.669

Metric	Value
MSE (Mean Squared Error)	24.29
RMSE (Root Mean Squared Error)	4.93

- **R² Score of 0.669** means that the model explains ~66.9% of the variance in the target variable.
 - **Residuals** were normally distributed, indicating the model is fairly unbiased.
 - **Predicted vs Actual Plot** showed good linear alignment but also some variance for higher-priced houses.
-

Conclusion

- The linear regression model performs reasonably well on this dataset.
 - It can be improved by:
 - Feature engineering (interaction terms or polynomial features)
 - Using more robust models like Ridge, Lasso, or Random Forest
 - Applying cross-validation to reduce overfitting
-

TASK 2 - Model Fit and Residual Plots

Figure 1: Model Fit - Actual vs Predicted Prices

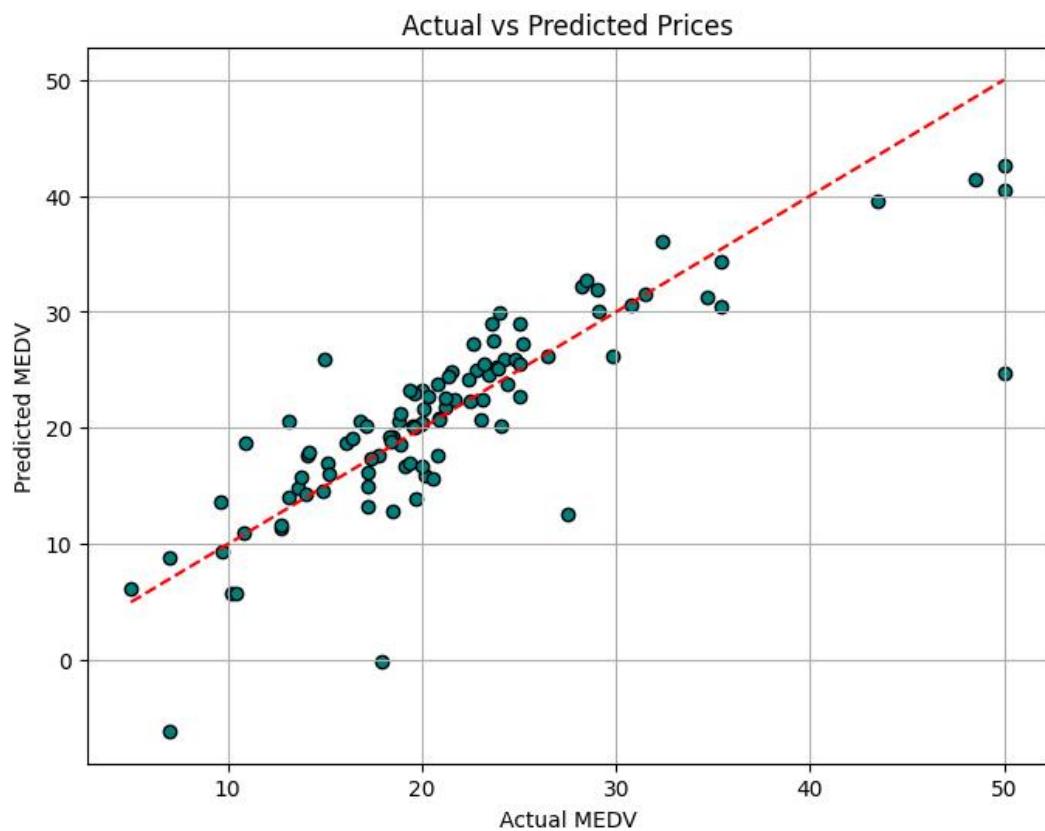


Figure 2: Residual Plot - Distribution of Residuals

